

Assignment Project Exam Help

XJCO3221 Parallel Computation

<https://eduassistpro.github.io>

University of Leeds

Add WeChat edu_assist_pro

Lecture 5: Loop parallelism and data

In Lecture 3 we saw how two problems could be parallelised:

- 1 **Vector addition**, where two one-dimensional arrays were added together element-by-element.

2

<https://eduassistpro.github.io>

However, neither of these problems have any

- 1 Each vector element was calculated **independently** of the others.
- 2 Each pixel colour was calculated **independently** of the others.

Assignment Project Exam Help

Today we will look how to parallelise problems that **do** have data dependencies.

- Can lead to **data races** on shared memory systems.



- <https://eduassistpro.github.io>

Add WeChat edu_assist_pro

We will then look at three examples of how their dependencies can be resolved.

Example of a data race

Consider the following pseudocode¹ for two concurrent threads, where each thread accesses the same variable x . $x=0$ at the start of the code segment.

$x = a;$

What value does x take at the end?

¹From §2.6 of McCool *et al.*, *Structured parallel programming* (Morgan-Kaufman, 2012).

The result may differ each time the program is executed.

- An example of non-determinism [Lecture 5]

The s
othe

- Cannot predict which thread is **la**
- The OS may **suspend** threads to run
(*pre-emptive multitasking*).
- The instructions may become **interleaved**.

Interleaved instructions (example)

Recall $x=0$ initially.

```
1  a  = x;    // Thread 0: a now 0
2  a  = 1;    // Thread 0: a now 1
3  b  = x;    // Thread 1: b now 0
4  x  = a;    // Thread 0: x now 1
5  b  += 2;   // Thread 1: b now 2
6  x  = b;    // Thread 1: x now 2
```

Add WeChat edu_assist_pro

In this example, $x=2$ at the end.

- Possible to get $x=1$ or $x=3$ by different interleaving of instructions (*check left as exercise*).

This is known as a **data race** or a **race condition**.

- Result of calculation depends on which thread reaches its instructions first.



<https://eduassistpro.github.io>

Only an issue for **shared memory**.

- If each thread had its own x , there would be no issue.
- In this example, if each thread had its own x , thread 0 would have $x=1$ and $x=2$ for thread 1 at the end, regardless of any interleaving.

Read-only does not lead to a data race

For a race condition to arise, **at least one** thread must **write** to `x`.

- No race if all threads just **read** `x`.

There
read

<https://eduassistpro.github.io>

Thread 0: _____

```
a = 2;  
a += 1.0;
```

Add WeChat edu_assist_pro

For this example, `x=0` (and `a=1` and `b=2`) at the end.

Have assumed each thread executes its instructions in order.

- i.e. have assumed **sequential consistency**.

Mod
perfo

- e.g. bring forward memory accesses, *co* *tc*.
- The result is the same **in serial**.
- However, **multithreading** can conf
- Can lead to unexpected results!

The volatile keyword

You may read that the way to solve this is to declare variables as `volatile` (in C/C++). However, this is only partially correct¹.

- `volatile` is for special memory such as that read/written by external device (*memory-mapped I/O*).

<https://eduassistpro.github.io>

If this might be an issue, should use features concurrent programming.

- e.g. memory fences, `std::atomic`
- Will come to atomics next lecture and Lecture 18.

¹S. Meyers, *Effective modern C++* (O'Reilly, 2015).

Often we are required to parallelise **loops**.

- Known as **loop parallelism**

- If there are **data dependencies**, may have implicit **data**

Then

- How to remove a dependency depends on co
- Consider extra resources or loops to reach th

For the remainder of this lecture, will give examples of loops with data dependencies, and how to overcome them.

Example 1: Redundant variable

Consider the following serial code:

```
1 float temp, a[n], b[n], c[n];  
2 ... // Initialise arrays b and c  
3  
4 int i;  
5 for( i=0;  
6 {  
7     te  
8     a[i] = temp;  
9 }
```

Here, temp is being used as a temporary variable

- Sometimes useful to make (more complex) code easier to read.

Need to make temp a private (or local) variable:

```
1 #pragma omp parallel for
2 for( i=0; i<n; i++ )
3 {
4     float temp = 0.5f*( b[i] + c[i] );
5     a[i] = temp;
6 }
```

Can also

```
1 #pragma
2 for( i=0; i<n; i++ )
3 {
4     temp = 0.5f*( b[i] + c[i] );
5     a[i] = temp;
6 }
```

cf. the inner loop counter in Lecture 3's Mandelbrot set example.

Example 2: Shift dependency

Code on Minerva: `shiftDependency.c`

Consider a shift dependency:

```
1 float a[n];  
2 ... // initialise array a  
3  
4 int i;  
5 for( i=0  
6     a[
```

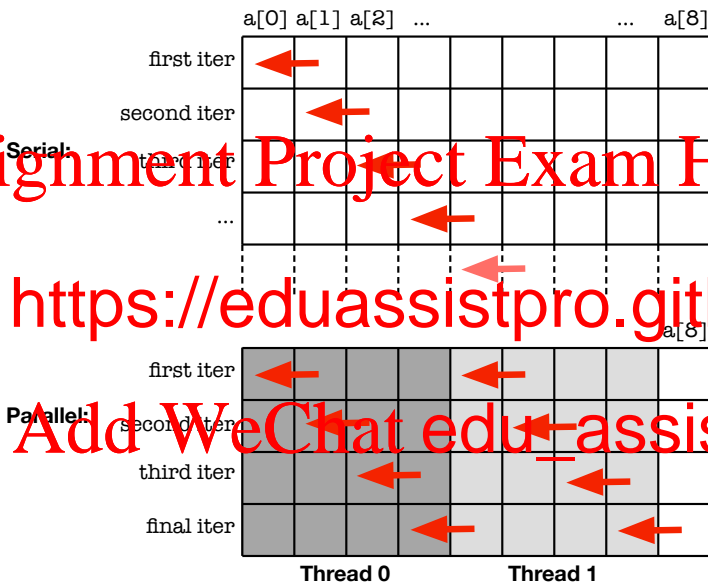
Naive parallelisation does not *quite* work

```
1 #pragma omp parallel for  
2 for( i=0; i<n-1; i++ )  
3     a[i] = a[i+1];
```

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



A solution here is to **copy** the array **a** **before** the loop:

```
1 float atemp[n];
2
3 #pragma omp parallel for
4 for( i=1; i<n; i++)
5     atemp[i] = a[i];
6
7 #pragma omp parallel for
8 for( i=0; i<n; i++)
9     a[i] = atemp[i];
```

This comes at the expense of **additional**

- **Memory** for the array **atemp**.
- **CPU time** to copy **a** to **atemp**.

Examples of **parallel overheads**.

Example 3: Red-black Gauss-Seidel

Code on Minerva: `redBlackGaussSeidel.c`

Finally, consider this:

```
1 for( i=1; i<n-1; i++ )  
2   a[i] = 0.5f * ( a[i+1] + a[i-1] );
```

Each

- Can be used to **smooth** vector a , transformation: ("blurring") with a 2D re
- Also arises in numerical computation - the **equation** solved using the **Gauss-Seidel** method.

We **could** make a copy atemp as before.

- In numerical computation, this is the **Jacobi method**.

Assignment Project Exam Help

However, this is undesirable in some situations:



<https://eduassistpro.github.io>

Instead we consider a **modified serial**
amenable to parallelisation.

- Known as **red-black Gauss-Seidel**
to red and black squares on a chessboard (in 2D).


Update **even** elements first, then **odd** elements¹:

```
1 int redBlack;
2 for( redBlack=0; redBlack<2; redBlack++ )
3     for( i=1; i<n-1; i++)
4         if( i%2 == redBlack )
5             {
6
7
```

The o

- When redBlack=0, only the elements o **n** are updated.
- Similarly, only the **odd** are updated

¹Recall $i\%2$ gives the remainder of division by 2, so e.g. $i\%2==0$ if i is even.



Assignment Project Exam Help


redBlack

<https://eduassistpro.github.io>

Update even
('red') sites:

redBlack=1

Add WeChat edu_assist_pro



Update odd
('black') sites:

Note that for each loop, the calculations are **now independent**.

- We have **removed the dependency**, albeit by slightly changing the serial algorithm.

Assignment Project Exam Help

Now clear how to parallelise:

```
1 for( red
2     #p
3     fo
4     if( i%2 == redBlack )
5     {
6         a[i] = 0.5f * ( a[i-1] + a[i+1] );
7     }
```

There are no dependencies **within** the i-loop, because $a[i-1]$ and $a[i+1]$ were/will be updated in the other redBlack loop.

The 'best' parallel algorithm?

Notice that we changed Gauss-Seidel to the red-black variant to make it more efficient in parallel.

- For Gauss-Seidel, this is the standard parallel variant.

Perf

①

- ② Parallel scaling is good (for large arrays).

As a general observation, the 'best' parallel algorithm need **not** be directly related to the 'best' serial algorithm.

Today we have seen how multiple threads with at least one writing to shared memory can lead to data races.

- Outcome is not predictable (*non-deterministic*).



m loops.

- <https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Next time we will look at a way to **sy**
parallel program and apply it to a linked list.