

# Assignment Project Exam Help

<https://eduassistpro.github.io>

University of Leeds  
Add WeChat edu\_assist\_pr

Lecture 6: Critical regions and a

## Previous lecture

# Assignment Project Exam Help

In the last lecture we looked at **data races** and **loop parallelism**:



<https://eduassistpro.github.io>

- Leads to **non-deterministic** beha

- Can arise in loops as **data depend**

- Often possible to remove these dependenc

the expense of increased **parallel overheads**.

## This lecture

# Assignment Project Exam Help

In this lecture we will look at an important concept in parallel programming: **synchronisation**.



ing

- <https://eduassistpro.github.io>
- Briefly mentioned in Lecture 4 in the context of **fork-join**.

Now we will focus on using synchronisation to avoid

- Define **critical regions** which can only be executed by one thread at a time.
- **Atomics**: Critical regions specialised to single operations.

We will say more about atomics in Lecture 18.

## Singly linked lists


Serial code on Minerva: `linkedList.c`

# Assignment Project Exam Help

Linked lists are a form of **container** in which each item is linked together in a **chain**:



<https://eduassistpro.github.io>



Note this is a *singly* linked list - a **doubly** linked list has arrows 'both ways,' *i.e.* a field `item_t *prev;`.

## List storage

For today's example, the data for each item is just a single integer.

# Assignment Project Exam Help

To link into a list, convenient to use a struct in C:

```
1  typedef
2  int data_t;
3  struct
4  } item_t;
5
6  item_t *head = NULL; // First item
```

New items are added using addToList:

```
1  for( i=0; i<numAdd; i++ )
2  addToList( i );
```

## Implementation of addToList()

```
1 void addToList( int data )
2 {
3     item_t *newItem = (item_t) malloc(sizeof(item_t));
4     newItem->data = data;
5     ne
6
7     // Find
8     if h
9     {
10         head = newItem;
11     }
12     else // Find end of list and add to it.
13     {
14         item_t *tail = head;
15         while( tail->next != NULL ) tail = tail->next;
16         tail->next = newItem;
17     }
18 }
```

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro

## Linked lists in parallel

# Assignment Project Exam Help

Suppose we want to add multiple items in parallel for speed.

The obvious thing is to create multiple threads, and have each thread

In Op

<https://eduassistpro.github.io>

```
1 #pragma omp parallel for
2 for( i=0; i<numAdd; i++ )
3     addToList( i )
```

Add WeChat edu\_assist\_pr

- Multiple threads created at start of loop (‘
- Each thread calls addToList some fraction of numAdd times.
- Check with printList() after the loop is complete (‘join’).

## Failure of `addToList()` when called in parallel

The items will not be added in the same order as in serial, but this may not be a problem as long as they are *somewhere* on the list.

Som

- <https://eduassistpro.github.io>

Also, the memory allocated for lost items is never re

- This **memory leak** could cause prob  
was run for a long time (e.g. serv

The implementation of `addToList()` does not work in a **multithreaded context**. We say it is not **thread safe**.



## Thread safety

A routine, library class, etc. is called **thread safe** if it works 'as advertised' in a **multithreaded context**.

If an A

<https://eduassistpro.github.io>

Note that being thread safe does **not** **efficient** in parallel.

- May have used a 'lazy' method to make a routine but very slow.
- In this case may need to find an alternative that *does* scale in parallel, or develop your own solution.

## The need for synchronisation

# Assignment Project Exam Help

This g

①

<https://eduassistpro.github.io>

②

When traversing the list, multiple threads may reach the tail at the same time. Again, only one is added.

③

A thread can read 'tail->next'!  
another thread sets 'tail->next'

If also *removing* (or 'popping') items, similar considerations would apply, although things would be more complicated.

## Critical regions

# Assignment Project Exam Help

It is not easy to remove the data dependencies in this case.



<https://eduassistpro.github.io>

An alternative strategy is to ensure only one thread

**critical regions** of code at a time.

- Implemented in OpenMP as `#pragma`
- Called **lock synchronisation**, for reasons that will become clear next lecture.

```
#pragma omp critical
```

# Assignment Project Exam Help

Critical region

0

<https://eduassistpro.github.io>

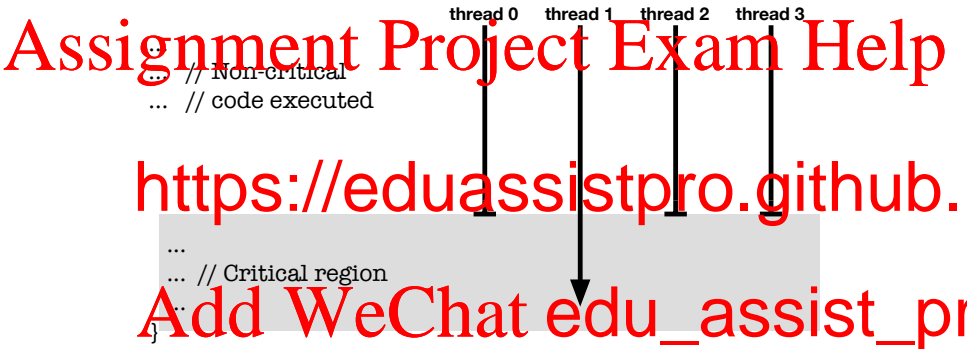
A **critical region** is defined in OpenMP as foll

```
1 #pragma omp critical
2 {
3     ... // Critical region
4 }
```

Add WeChat edu\_assist\_pro

The critical region is defined by the **scope**, *i.e.* from '{' to '}'.

## Example for 4 threads



- Thread 1 reaches the critical region **first**.
- **No other thread** can enter until it leaves.
- **Exactly one** thread may then enter.

## Performance

# Assignment Project Exam Help

There can be a significant performance penalty for critical regions:

①

<https://eduassistpro.github.io>

②

**load balancing** [*cf. Lecture 13*].

③

The scheduler may suspend idle threads (not necessarily yours!). Suspension and resumption penalties.

## Serialisation

Since only one thread can be in a critical region at any time, the critical code is **executed in serial**.

Am

<https://eduassistpro.github.io>

- Maximum speed-up  $S$  in terms of the
- By serialising regions of code we are
- The maximum speed-up  $S$  is **red**  $f$ .  
law (*i.e.* strong scaling), which predicts  $S \leq 1/f$ .

It is therefore important to restrict the **number** and **size** of critical regions to ensure reasonable parallel performance.

## First attempt: Serialise calls to addToList()

# Assignment Project Exam Help

```
1 #pragma omp parallel for
2 for( i=0; i<numAdd; i++ )
3     #pragma omp critical
4     {
5
6 }
```

<https://eduassistpro.github.io>

This works, but parallel performance is poor.

- Essentially the whole loop has been serialis
- Would be better off leaving it in serial,

```
1 for( i=0; i<numAdd; i++ )
2     addToList( i );
```



## Attempt 2: Serialise list traversal

Note that the start of `addToList()` has no data dependences

```
1 item_t *newItem = (item_t*) malloc(sizeof(item_t));  
2 newItem->data = data;  
3 newItem->next = NULL;
```

- <https://eduassistpro.github.io>
- Each thread will create its own item **independently of other threads**.
- Each value of loop counter `i` will have a unique value of `data`.

This is the behaviour we want! (so far...)

The data dependencies in the remainder of `addToList()` can be removed by placing this portion in a critical region:

```
1 #pragma omp critical
2 {
3     if h
4     {
5
6     }
7     else
8     {
9         item *tail = head;
10        while (tail->next != NULL) tail = tail->next;
11        tail->next = newItem;
12    }
13 }
```

Performance *slightly* improved compared to the previous attempt.

## Making routines thread safe

# Assignment Project Exam Help

Note that

is not thread safe:

①

②

③

④

<https://eduassistpro.github.io>

**Reduce** size and/or number of critical regions to achieve the performance achieved.

If further scaling benefits required, may need to change algorithm completely.

Add WeChat edu\_assist\_pro

## Atomics

# Assignment Project Exam Help

Often the data dependency is only a single arithmetic operation.

For instance, counting the number of items in an array of size  $n$  that o

```
1 int count = 0;
2 for( i=0; i<n; i++)
3 {
4     ...
5     if( condition[i] ) count++;
6 }
```

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pr

The command `count++` is a **data race**:

- Two threads may read the **old** value of `count` simultaneously.
- New `count` may not be the old value  $+2$  [*cf. Lecture 5*].

## Critical region

# Assignment Project Exam Help

Using a critical region will work ...

```
1 int count = 0;
2 #prag
3 for( i=0
4 {
5     ..
6     if( condition[i] )
7         #pragma omp critical
8         {
9             count++;
10        }
11 }
```

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pr

... but has the usual overheads of a critical region.

## Atomic instructions

Assignment Project Exam Help

Because this is a common situation, most compilers/hardware can perform the necessary synchronisation efficiently.



<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pr

Can think of atomics as a special type of critical region for single arithmetic operations that exploits hardware

GPUs also support atomic instructions; we will look at these more closely in Lecture 18.

## Atomics in OpenMP

# Assignment Project Exam Help

Atomics are implemented in OpenMP as follows:

```
1 int count = 0;
2 #prag
3 for( i=0
4 {
5     ..
6     if( condition[i] )
7     #pragma omp atomic
8         count++;
9 }
```

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pr

Note there is no scope ('{'...'') after `#pragma omp atomic`.

- Only works on single instructions.

## Summary and next lecture

# Assignment Project Exam Help

Today we have looked at using **critical regions**:

- 
- 
- 

<https://eduassistpro.github.io>

atomics.

## Add WeChat edu\_assist\_pr

Next lecture we will look in more detail at how this synchronisation is achieved.