

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat University of Leeds edu_assist_pr

Lecture 7: Lock and mute

Previous lecture

Assignment Project Exam Help

In the last lecture we saw how critical regions of code could be serialised

- <https://eduassistpro.github.io>
- Avoids **data races**.
- Can incur a significant **performance** overhead.
- Implemented in OpenMP as `#pragma omp atomic`
- Single arithmetic instructions can be optimised by using **atomic** instructions (`#pragma omp atomic`).

Today's lecture

For today's final lecture on shared memory parallelism, we will look at what is going on 'behind the scenes'.

- Thread coordination performed using **locks**, sometimes known

- <https://eduassistpro.github.io>
- ess.

- However, multiple locks can give rise to

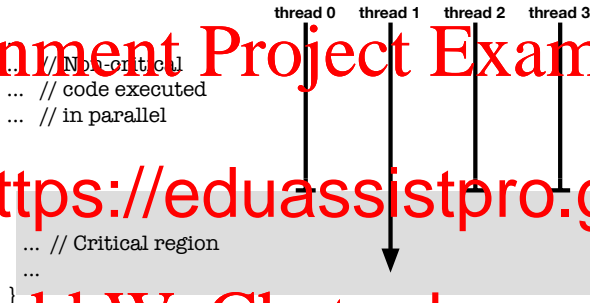
This lecture is largely theoretical¹ a
courseworks, but the material may appear in the exam.

¹There **are** code examples for this lecture, and a question on Worksheet 1.

Recap: Critical regions

Assignment Project Exam Help

<https://eduassistpro.github.io>



Add WeChat edu_assist_pr

- Instructions before `#pragma omp cri`
concurrently (e.g. if in a parallel loop).
- Instructions in the scope (`{` to `}`) only executed by **one thread at a time**.
- Other threads blocked from entering; they are **idle**.

Thread coordination with locks

Assignment Project Exam Help

This synchronisation is performed using a **lock**:



<https://eduassistpro.github.io>

- Also known as **acquiring the lock**.

- This thread is said to be the lock's



No other threads can enter the region (it's **acquired**) until it becomes unlocked.



The owning thread **unlocks** (or **releases**) it when leaving the region, allowing another thread to take over ownership.

Critical region using a lock

OpenMP:

```
1 // Multiple threads exec-  
2 // uting co  
3 // (e.g. pa  
4  
5 #pragma  
6 {  
7     ...  
8     ... // Critical code.  
9     ...  
10 }
```

Lock pseudocode:

```
1 // All threads access a  
2  
3  
4  
5  
6  
7 ..  
8 .. // C  
9 ..  
10 re
```

`regionLock.lock()` does not return until the thread has **acquired** the lock; it is said to be **blocking**.

Implementations of locks

Assignment Project Exam Help

Most parallel APIs support **locks**, although they are sometimes called **mutexes** as they control m**utual** e**xclusion**:

- <https://eduassistpro.github.io> (cks).
- `pthread_mutex_t` in the pthreads library (C/C++).

Add WeChat edu_assist_pr

When implemented as classes, they are typically—

- The user does not have access to instance variables or details of the implementation.

Locks in OpenMP

OpenMP runtime library also supports locks:

```
1 #include <omp.h>
2
3 // Init
4 omp_lock_t regionLock;
5
6 ... // (in parallel).
7 omp_set_lock(&regionLock); // LOCK.
8 ... // (critical code)
9 omp_unset_lock(&regionLock); // UNLOCK.
10 ...
11 // Deallocate the lock.
12 omp_destroy_lock(&regionLock);
```

You *could* implement your own critical region this way, although it is easier to use `#pragma omp critical`.

Programming locks

Assignment Project Exam Help

Note there is no explicit link between the lock region, lock and the critical region of code, or data structure, that it is trying to protect

It is doing an associated block of critical code, or data

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

This gives greater **flexibility**, but also great programming errors.

- Could use a struct or class to keep the lock with the data it is protecting, with the lock private/protected.

Lock mistakes (1): Forgetting to lock()

```
1 lock_t regionLock;  
2 ...  
3 //regionLock.lock(); // Forgot to lock()!  
4 ...  
5 ... // Criti  
6 ...  
7 region
```

This is precisely the situation we were trying to avoid

- All threads enter the critical region
- **Race conditions** become a possibility

unlock() will have no effect, except possibly a small performance overhead¹.

¹Generally, this depends on the API: In C++11, attempting to unlock a `std::mutex` that is **not** locked leads to undefined behaviour.

Lock mistakes (2): Forgetting to unlock()

Assignment Project Exam Help

```
1 lock(&regionLock;  
2 ...  
3 regio  
4 ...  
5 ... // Bruti  
6 ...  
7 //regionLock.unlock(); // Forgot to unlock!
```

<https://eduassistpro.github.io>

- The first thread **exclusively** enters
- It never **releases** the lock.
- Therefore no other thread can **acquire** the lock.
- **All other threads remain idle at lock().**

Add WeChat [edu_assist_pro](https://eduassistpro.github.io)

RAII = Resource Acquisition Is Initialisation.

Assignment Project Exam Help

This second mistake is easier to make than it seems:

- The critical code may throw an **exception** (C++/Java).



ck().

<https://eduassistpro.github.io>

May be

leave their scope.

- If defined at start of a routine, automatically released at end of routine **however it reached there**

- e.g. `std::lock_guard<std::mutex>` in C++11.

This mechanism is generally known as RAII, for Resource Acquisition Is Initialisation.

Multiple locks

Code on Minerva: `multipleLockCopy.c`

Assignment Project Exam Help

Suppose we want to copy randomly selected elements of an array of size N to another randomly-selected element.

```
1 #pragma omp parallel for
2 for( n=0; n<N; n++)
3 {
4     i = rand() % N;
5     j = rand() % N;
6
7     omp_set_lock( &entireLock );    // Lock.
8     data[j] = data[i];              // Safe copy.
9     omp_unset_lock( &entireLock );  // Unlock.
10 }
```

Multiple locks for memory access

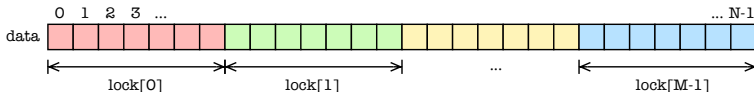
This works, but is very inefficient.

- Only one thread can access the array data at a time.



Better to use **multiple locks** spanning t

- Different threads can write to different regions **simultaneously**.
- Less **idle time** spent waiting for a lock to be released.



Using multiple locks is measurably faster (*try the code*):

```
1 omp_lock_t partialLocks[M];  
2  
3 // Initialise M locks near start of code.  
4 ...  
5 // Ident  
6 int lock = M*  
7 omp_s  
8 data[  
9 omp_unset_lock( &partialLocks[lock] );  
10 ...  
11 // Destroy all locks at end of code.
```

Note we only lock for the **write** to element j

- Recall that just reading does **not** invoke a data race.

Multiple locks for swapping

Code on Minerva: `multipleLockSwap.c`

Assignment Project Exam Help

Suppose now we want to swap elements i and j .

- Want to protect **each write** during the swap.

If access
would

<https://eduassistpro.github.io>

```
1 omp_set_lock( &entireLock );
2
3 // Writes to both data[i] and data[j].
4 float temp = data[i];
5 data[i] = data[j];
6 data[j] = temp;
7
8 omp_unset_lock( &entireLock );
```

However, performance would again be poor.

Multiple locks for swapping

We might think of using two locks, one for each region of the array being written to:

```
1 int lock_i = M*i/N;  
2 int lock  
3  
4 omp_s  
5 omp_s  
6  
7 float temp = data[i];  
8 data[i] = data[j];  
9 data[j] = temp;  
10  
11 omp_unset_lock( &partialLocks[lock_i] );  
12 omp_unset_lock( &partialLocks[lock_j] );
```

Try this out!

Why does this fail? Deadlock

Assignment Project Exam Help

Suppose one thread tries to lock `lock_i` then `lock_j` and **simultaneously** another tries to lock `lock_j` then `lock_i`.



<https://eduassistpro.github.io>

Sinc

release the lock they own. **They will bot**

Add WeChat edu_assist_pr

Threads waiting for synchronisation events th
known as **deadlock**.

The 'forgetting to `unlock()`' example earlier is also **deadlock**.

Nested critical regions

Code on Minerva: `nestedCriticalRegion.c`

Assignment Project Exam Help

Another problem is when `lock_i == lock_j`. A simpler example where this occurs is for **nested critical regions**:

```
1 // Outer c
2 omp_s
3
4 // Inner c
5 omp_set_lock( &lock );
6 ...
7 omp_unset_lock( &lock ); // End of inner region
8
9 omp_unset_lock( &lock ); // End of outer region
```

In OpenMP, this will also **deadlock**.

- A thread that **owns** a lock cannot **re-acquire** the lock.

Nested #pragma omp critical

Assignment Project Exam Help

OpenMP does not allow nested critical regions:

```
1 #pragma omp critical
2 {
3     ..
4     #p
5     {
6         ...
7     }
8 }
```

... will **not** compile.

- The **same** lock is being used by **both** critical sections.
- The same problem as in the previous slide.

Named critical regions

This can be resolved by using **named** critical regions:

```
1 #pragma omp critical (OUTER)
2 {
3     ..
4     #p
5     {
6
7     }
8 }
```

- OUTER and INNER are user-defined labels
- Each unique label corresponds to a unique lock.
- You are implicitly using a different lock for each critical region, so no thread tries to re-acquire a lock it already owns.

Reacquiring locks

OpenMP code **deadlocks** if a thread tries to reacquire a lock it already owns.



This is



necessary.

Not all parallel/concurrent APIs impose the same
need to check the documentation!

- e.g. For C++11's `std::mutex`, the behaviour is undefined.
- Should also check documentation if attempting to `unlock` a lock that was **not** acquired.

Summary of shared memory systems

Assignment Project Exam Help

Lecture	Content	Key points
	Arch	ernel
	Data par	Finishing
4	Theory	Amdahl's law (strong scaling); Gustafs
5	Data races	Loop parallelism, data dependencies
6	Critical regions	Thread coordination; thread safety; serialisation; atomics.
7	Locks/mutexes	Performance costs for locks; dead-lock; named critical regions.

Next lecture

Assignment Project Exam Help

Next time we will start to look at **distributed memory systems**:



<https://eduassistpro.github.io>

Not surprisingly, data races are *not* an interesting aspect we have covered are:

- Non-determinism, scaling, deadlock, d
parallelism, . . .

Add WeChat edu_assist_pro