

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat University of Leeds edu_assist_pr

Lecture 18: Atomic opera

Previous lectures

Assignment Project Exam Help

Whenever multiple processing units had read-write access to the same memory location, there are potential **data races**:



<https://eduassistpro.github.io>



Single instructions can also be made perform

[Lecture 6].

Add WeChat edu_assist_pro

For instance, in OpenMP an atomic instruction looks like:

```
1 #pragma omp atomic
2 count++;
```

Today's lecture

Assignment Project Exam Help

GPUs also have memory accessible by multiple work items / threads:



<https://eduassistpro.github.io>

Then

synchronisation.

Add WeChat edu_assist_pr

Today we will see how GPUs support similar way to a shared memory CPU.

- Also consider an atomic **compare and swap**.

Atomic operations

Assignment Project Exam Help

Definition

ut

<https://eduassistpro.github.io>

- Usually restricted to simple **arith** on, subtraction etc.)
- Implemented by a combination of compil
- Typically a (much) **smaller performance penalty** than using locks/mutexes etc.

Add WeChat edu_assist_pr

Load, compute, and store

Consider the following line:

```
1 x -= 2;
```

Even though

①

②

③

<https://eduassistpro.github.io>

Stores the updated value.

Add WeChat edu_assist_pro

Two or more processing units might interfere with each other, resulting in a different result to the **serial** equivalent.

This could not happen if the operation was **atomic**.

Example

Assignment Project Exam Help

Suppose $x=10$ initially, and two processing units A and B both subtract 2 from x . Depending on the scheduler, this may happen:

①

②

③

④ A stores 8 to memory.

⑤

<https://eduassistpro.github.io>

⑥

B stores 8 to memory.

The result is $x = 8$, rather than $x = 6$ as expected.

Constructing a histogram on a GPU

Code on Minerva: `histogram.c`, `histogram.cl`, `helper.h`

Assignment Project Exam Help

Have an array of integers in the range 0 to `maxValue-1` inclusive; want the **histogram** showing the frequency of each value.

①

th

<https://eduassistpro.github.io>

- `device_data`, `device_hist` on the device.

②

Both initialised on the host and copied to the d

③

Build, initialise and enqueue a kernel to build the histogram.

- **One work item per data element**, e.g. `data[i]`.

④

Copy the histogram back to the host using `clEnqueueReadBuffer()`.

Kernel 1: Direct to global; no atomics

Assignment Project Exam Help

```
1 // kernel
2 void histogramNoAtomic(
3     --
4     --
5 {
6     int gid = get_global_id(0);
7
8     // Data value.
9     int val = device_data[gid];
10
11     // Check range before updating.
12     if( val >= 0 && val < maxValue )
13         device_hist[val]++;
14 }
15 }
```

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

Kernel 2: Direct to global; atomic.

Code fails because the update of `device_hist` is not atomic

```
1 device_hist[val]++;           // Load, compute, store.
```

Man

In Op

<https://eduassistpro.github.io>

```
1 __kernel
2 void histogramAtomic( ... )
3 {
4     ... // as before
5     if( val >= 0 && val < maxValue )
6         atomic_inc( &(device_hist[val]) );
7 }
```

This now works as expected.

Atomic operations in OpenCL

Assignment Project Exam Help

There are many atomic operations in OpenCL¹:

<code>atomic_inc</code> , <code>atomic_dec</code>	Increment, decrement.
<code>a</code> <code>-</code>	
<code>a</code>	
<code>atomic_and</code> , <code>atomic_or</code> ,	Bitwise
<code>atomic_xor</code>	
<code>atomic_fetchg</code> ,	Exchange
<code>atomic_cmpxchg</code>	compare and exchange

¹Similar in CUDA, i.e. `atomicAdd()`, `atomicInc()` etc.

Optimising with local memory

Assignment Project Exam Help

Having a **single** histogram in global memory is not efficient:

- Potential for **very many** work items attempting to access the

<https://eduassistpro.github.io>

More

in **local memory**, then update the global hist

- Fewer **competing** work items for the local memory
- Local memory is faster anyway.

Add WeChat edu_assist_pro

Aside: Could use a similar strategy for a *multi-threaded CPU* (i.e. each *thread* constructs its own histogram).

Kernel 3: Local histogram (1)

```
1 _kernel
2 void localHistogram(..., __local int *local_hist)
3 {
4     in
5
6
7
8     size = get_local_size(0);
9
10    // Clear the histogram
11    for(i=id; i<maxValue; i+=size)
12        local_hist[i] = 0;
13
14    // Ensure histogram fully initialised.
15    barrier(CLK_LOCAL_MEM_FENCE);
16
17    // (cont'd next slide).
```

Kernel 3: Local histogram (2)

```
1 // (from previous slide)
2
3 // Add to the local histogram.
4 in  v
5 if  v
6
7
8 // Ensure local histogram calculation complete
  before moving on.
9 barrier(CLK_LOCAL_MEM_FENCE);
10
11 // Atomic add the local histogram to the global one.
12 for( i=lid; i<maxValue; i+=size )
13     atomic_add( &(amp;device_hist[i]), local_hist[i] );
14 }
```

You should see a performance improvement using this method.

Could have had **one** work item in each group initialise and update the **entire** local histogram, e.g.:

```
1 if( gid==0 )  
2   for( i=0; i<maxValue; i++ )  
3
```

This <https://eduassistpro.github.io>

Instead use as many work items as possible:

```
1 for( i=gid; i<maxValue; i+=size )  
2   atomic_add( &device_hist[i], local_hist[i] );
```

- Each i in the range realised by **exactly** one work item.
- Spans full range even if $\text{size} < \text{maxValue}$.

Atomic exchange and compare-and-exchange

Assignment Project Exam Help

```
int atomic_exchg(int *p, int val).
```

① Sets old=*p.

②

③

<https://eduassistpro.github.io>

```
int atomic_cmpxchg(int *p, int cmp, int val):
```

① Sets old=*p

② Sets *p=val if *p==cmp, otherwise d

③ Returns old.

In both cases, p can be in local or global memory, and the data type can be int, unsigned int or float.

Uses of compare and exchange

Many low-level parallel frameworks provide similar functionality, sometimes referred to as CAS = Compare And Swap:

-

-

<https://eduassistpro.github.io>

Common uses include:

- 1

Implementing a lock or mutex

- 2

Lock-free implementations.

Examples below are for OpenCL, but just as relevant for CUDA and multi-core CPUs.

Spinlock

Suppose an `int` variable `lock` was accessible to multiple threads.

- `lock` takes **two** values, 0 and 1.

-

A sim

```
1 int lock;           // 0 or 1. Accessible by all threads.  
2 ...  
3 while (atomic_compare_exchange(&lock, 0, 1) != 1);
```

- Infinite while loop, until `lock==0`.
- Then sets `lock=1` and continues past line 3.
- Does all this **atomically**.

Why atomic?

Assignment Project Exam Help

Consider what could happen without atomicity:

```
1 while( lock==1 );  
2 lock = 1;  
3 ...
```

<https://eduassistpro.github.io>

- ➊ One thread / work item sees `lock==0` and continues to line 2.
- ➋ A second thread **also** sees `lock=`
before the first thread sets `lock`
- ➌ The first thread now sets `lock=1`
- ➍ The second thread **also** sets `lock=1`.
- ➎ **Both** continue to line 3!

Spinlocks vs. locks/mutexes:

- **Pro:** Spinlocks **faster**, as do not put the thread to sleep.
- **Con:** Spinlocks waste CPU/GPU cycles.

Spin

Note



- May seem this can be used to **sy** ps.

But recall the warning from Lecture 17:

Cannot guarantee all work groups are active on the device **at the same time** (as some may be queued), so this not a **robust** synchronisation mechanism.

Lock-free data structures

Atomic compare and exchange can also be used to implement thread-safe access to data structures without requiring locks, and th

<https://eduassistpro.github.io>

Example: Prepending an item to a singly

- Need to ensure old and new head
- Use `atomic_cmpxchg()` in an infinite loop¹.

¹McCool *et al.*, *Structured parallel programming* (Morgan-Kaufman, 2012).

Basic idea of a lock-free data linked list

```
1 struct node { ... }; // Some data structure
2 node *head; // head of list
3 void prependToList( node *a ) // 'a' becomes head.
4 {
5     wh
6     {
7         node *b = head;
8
9         // Link to the node being added
10        a->next = b;
11
12        // Only update head if not just changed by another
13        // work item/thread; else try again from line 6.
14        if( atomic_cmpxchg(head,b,a)==b ) break;
15    }
16 }
17 }
```

```
if( atomic_cmpxchg(head,b,a)==b ) break;
```

Assignment Project Exam Help

If only a single thread was involved:

①

②

③

<https://eduassistpro.github.io> or a loop.

This is the expected behaviour.

However, in a multi-threaded context:

Add WeChat [edu_assist_pr](#)

①

Another thread may change *head b

②

Since *head!=b, will **not** change it.

③

Will return some value !=b, so will try again.

Summary and next lecture

Assignment Project Exam Help

This lecture we have revisited **atomic operations** with an
emp

- <https://eduassistpro.github.io>
 - Atomic **compare and exchange** can be used to implement a
spinlock, lock-free data structures,
- Add WeChat edu_assist_pr

The next lecture is the last on GPU programming when we will
look at events and **task parallelism**.