

Assignment Project Exam Help

<https://eduassistpro.github.io>

Peter Jimack

Add WeChat University of Leeds edu_assist_pr

Lecture 20: Summary

Previous lectures

Assignment Project Exam Help

This module has been structured to focus on one parallel archi

①

②

<https://eduassistpro.github.io>

③

General purpose GPU (with OpenCL); Lectures 14-19.

Add WeChat edu_assist_pr

The practical elements of the module (*worksheets and courseworks*) also followed this structure.

- Some *mentions* were out-of-order, e.g. OpenMP barriers mentioned in Lectures 11 and 17.

Today's lecture

Assignment Project Exam Help

Want to understand general parallel programming concepts that transcend particular architectures.

- <https://eduassistpro.github.io>
- e.

In this final lecture we will summarise all of the material
parallel concept rather than by architecture.

- Easier to see the commonalities.

At the end I'll also spend a few minutes talking about the **Final assessment** for the module!

Why parallel?

Lectures 1 and 4

Assignment Project Exam Help

Parallel hardware allows **simultaneous** computations.

- <https://eduassistpro.github.io>
-

could be e.g. time-sharing on a single c

Add WeChat edu_assist_pr

Want to attain good **scaling** - decrease i
time t_p for increasing number of **processing units** (*threads*,
processes etc.) p .

Measuring parallel performance

Lecture 4

Assignment Project Exam Help

Two useful metrics for parallel performance are the **speed-up** S and **efficiency** E

<https://eduassistpro.github.io>

Achieving $S = p$ (i.e. $E = 1$) usually regar
difficult to achieve due to various over

- Synchronisation, load balancing, com
calculations, ...
- **Super-linear scaling** $S > p$ possible (but rare) due to
memory cache.

Laws for maximum parallel performance

Lectures 4, 19

Assignment Project Exam Help

Strong scaling is when the problem size n is fixed.

- $\frac{1}{\text{span}}$

- <https://eduassistpro.github.io>

Weak scaling allows n to increase with

- Related to the **Gustafson-Barsis I**

Add WeChat edu_assist_pro

The **work-span model** provides another maximum S from **task-graphs** [Lecture 19].

- $S \leq (\text{work})/(\text{span})$, with **work** and **span** determined from the task graph.


Loop parallelism and data dependencies

Lectures 3, 5, 9, 15

Assignment Project Exam Help

- Initially looked at parallelising loops.
- If there are no **data dependencies**, is **data parallel** or a **map**.

- <https://eduassistpro.github.io>



Synchronisation

Lectures 7, 9, 11, 17

Assignment Project Exam Help

All but the simplest parallel problems require **synchronisation** between

- <https://eduassistpro.github.io>
- **attribution.**

Can lead to reduced **performance**.

- e.g. the extra operations required to achieve **Add WeChat edu_assist_pr**

Can also lead to **deadlock** [Lectures 7, 9].

- When one or more processing units wait for a synchronisation even that never occurs.

Load balancing and task parallelism

Lectures 13 and 19

Assignment Project Exam Help

Idle time is an example of poor **load balancing** [Lecture 13].

-
- <https://eduassistpro.github.io>

Can improve load balancing by using a **dynamic** scheduler. **tasks** are sent to processing units and become idle.

- This is an example of **dynamic** scheduling; can also be **static**.

Assignment Project Exam Help

Parallelising by **tasks** rather than loops is known as **task parallelism** [Lecture 19].



<https://eduassistpro.github.io>

In gen

- Can represent as a **task graph**, with **tasks** and **directed edges** denoting the dependencies.
- For tasks that take the same time, can define the **he** total number of tasks, and the **span** as the length of the critical path.

Data reorganisation

Lecture 10

Assignment Project Exam Help

Parallel data reorganisation can be indexed by **read** locations
(`gather`)

In shared memory systems, <https://eduassistpro.github.io> is an example of a **data race** (see later).

In distributed memory systems, can exploit **communication methods** that are usual

- One-to-many, many-to-one (also many-to-many).
- e.g. broadcasting, scattering and gathering.

Parallel hardware

Lectures 2, 8, 14, 16

Assignment Project Exam Help

Modern HPC clusters are increasingly using all three architectures:



re

- <https://eduassistpro.github.io>

Most multi-core CPUs usually have **m**
cache coherency and **false sharing** [

Network connectivity affects communi
hypercube often used [Lecture 8].

GPU's most suited for **data parallel problems** and have multiple
types of **memory** [Lectures 14, 16].

Data races / race conditions

Lectures 5, 6, 18

Assignment Project Exam Help

A **data race** potentially arises when two or more processing units read t

- <https://eduassistpro.github.io>
-

Can control using critical regions

Add WeChat edu_assist_pr

- Exclusive access by a single processing unit
- Simple critical regions can be implemented more efficiently (i.e. by compiler and hardware) as **atomics**.

Lower level control

Lectures 7, 18

Assignment Project Exam Help

At a lo

mut

- <https://eduassistpro.github.io>
- Improper use of multiple locks can result in **deadlock**.

Add WeChat edu_assist_pr

At an even lower level, locks can be implemented using **exchange** and **atomic compare-and-exchange** [Lecture 18].

- **Lock-free data structures** are desirable whenever possible.

Explicit communication

Lectures 9, 10, 12, 15, 19

Assignment Project Exam Help

If memory is distributed (in some sense), may need to use **explicit com**

- <https://eduassistpro.github.io>

- Between CPU and GPU, i.e. **host**

Add WeChat edu_assist_pr
Communication can be:

- **Blocking**: Returns once all resources safe to re-use.
- **Synchronous**: Does not complete until sender and receiver start their communication operations.

Latency hiding

Lectures 12, 19

Assignment Project Exam Help

Can improve performance by **overlapping** communication with computation

- <https://eduassistpro.github.io>

Often used with **domain partitioning**

[Lecture 12].

Add WeChat edu_assist_pro

Can also overlap host-device communication with a GPU [Lecture 19].

- Can also perform calculations on host and device simultaneously.

The end

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

I will now say a few words about the final assessment
this module...

The final assessment

Assignment Project Exam Help

This will be an online exam to be taken at 1600 (China time) on Tuesday 17th May 2022...

<https://eduassistpro.github.io>

on 17th May (unless you have been granted

- It is an “open book” exam, but you will receive no copying material directly from your lecturer
- Your overall module grade will be the sum of your scores on the 3 courseworks and your score (out of 50) on this final assessment
 - Hence the final assessment is worth 50% of the total marks

The final assessment (continued)

Assignment Project Exam Help

- The paper will consist of two questions worth 25 marks each

- <https://eduassistpro.github.io>

- The 2020 and 2021 papers were both “open book” but the students were given more time to complete

questions are longer than they will be this year

- The 2018 and 2019 papers were both 2 hour they were “closed book” – so some section questions
- See the announcement on Minerva for some further advice

Your revision questions

Assignment Project Exam Help

- The final Zoom session will be held at 1730 on Tuesday 10th May

<https://eduassistpro.github.io>

questions (and my responses)

- There will also be a Zoom session at 1730 today if you already have any questions for me!
- I will also answer any questions that you post on XJCO3221 Discussion Boards on Minerva prior to 1600 on 17th May
 - I will NOT answer any questions after the start of the final assessment!